

Big Data Analytics

– Chapter 7: Frequent Itemset Mining –

Prof. Dr. Emmanuel Müller
Chair of Data Science and Data Engineering

Content Overview

- Introduction
 - Transaction databases, market basket data analysis
- Mining Frequent Itemsets
 - Apriori algorithm, hash trees, FP-tree
- Simple Association Rules
 - Basic notions, rule generation, interestingness measures
- Summarizing Frequent Itemsets
 - Maximal, closed, non-derivable itemsets
- Summary

What is Frequent Itemset Mining?

Finding frequent patterns, associations, correlations, or causal structures among sets of items or objects in transaction databases, relational databases, and other information repositories.

- Given:

- A set of items $I = \{i_1, i_2, \dots, i_m\}$
- A database of transactions D , where a transaction $T \subseteq I$ is a set of items

- Task 1: find all subsets of items that occur together in many transactions.

e.g.: 85% of transactions contain the itemset {milk, bread, butter}

- Task 2: find all rules that correlate the presence of one set of items with that of another set of items in the transaction database.

e.g.: 98% of people buying tires and auto accessories also get automotive service done

- Applications:

Basket data analysis, cross-marketing, catalog design, clustering, classification, recommendation systems, etc.

Example: Basket Data Analysis

- Transaction database

$T = \{ \{ \text{butter, bread, milk, sugar} \};$
 $\{ \text{butter, flour, milk, sugar} \};$
 $\{ \text{butter, eggs, milk, salt} \};$
 $\{ \text{eggs} \};$
 $\{ \text{butter, flour, milk, salt, sugar} \} \}$

- Question of interest:

- Which items are bought together frequently?

- Applications

- Improved store layout; Cross marketing; Focused attached mailings / add-on sales

- Generalization of „association rules“ beyond marketing:

- $\text{buys}(x, \text{"diapers"}) \rightarrow \text{buys}(x, \text{"beers"})$
- $\text{major}(x, \text{"CS"}) \wedge \text{takes}(x, \text{"DB"}) \rightarrow \text{grade}(x, \text{"A"})$



Basic Notions I

- *Items* $I = \{i_1, i_2, \dots, i_m\}$: a set of literals (denoting items)
- *Itemset* X : Set of items $X \subseteq I$
- *Database* D :
 - Set of *transactions* T , each transaction is a set of items $T \subseteq I$
 - Transaction T *contains* an itemset X : $X \subseteq T$

The items in transactions and itemsets are sorted lexicographically:

- itemset $X = (x_1, x_2, \dots, x_k)$, where $x_1 \leq x_2 \leq \dots \leq x_k$

Length of an itemset: number of elements in the itemset

k-itemset: itemset of length k

Basic Notions II

- *Items* $I = \{i_1, i_2, \dots, i_m\}$: a set of literals (denoting items)
- *Itemset* X : Set of items $X \subseteq I$

The *support* of an itemset X is defined as: $support(X) = |\{T \in D | X \subseteq T\}|$

Frequent itemset: an itemset X is called frequent for database D iff it is contained in more than $minSup$ many transactions: $support(X) \geq minSup$

Goal 1:


Given a database D and a threshold $minSup$,
find all frequent itemsets $X \in Pot(I)$.

Basic Idea

Naïve Algorithm

- count the frequency of all possible subsets of I in the database
- *too expensive* since there are 2^m such itemsets for $|I| = m$ items

cardinality of power set



Basic Idea

The *Apriori* principle (anti-monotonicity):

Any non-empty subset of a frequent itemset is frequent, too!


$A \subseteq I$ with $\text{support}(A) \geq \text{minSup} \Rightarrow \forall A' \subset A \wedge A' \neq \emptyset: \text{support}(A') \geq \text{minSup}$

Any superset of a non-frequent itemset is non-frequent, too!

$A \subseteq I$ with $\text{support}(A) < \text{minSup} \Rightarrow \forall A' \supset A: \text{support}(A') < \text{minSup}$

Method based on the apriori principle

- First count the 1-itemsets, then the 2-itemsets, then the 3-itemsets, and so on
- When counting $(k+1)$ -itemsets, only consider those $(k+1)$ -itemsets where all subsets of length k have been determined as frequent in the previous step



The Apriori Algorithm

```
variable  $C_k$ : candidate itemsets of size  $k$ 
variable  $L_k$ : frequent itemsets of size  $k$ 
 $L_1 = \{\text{frequent items}\}$ 
for ( $k = 1$ ;  $L_k \neq \emptyset$ ;  $k++$ ) do begin
    produce candidates { // JOIN STEP: join  $L_k$  with itself to produce  $C_{k+1}$ 
                       // PRUNE STEP: discard  $(k+1)$ -itemsets from  $C_{k+1}$ 
                       that contain non-frequent  $k$ -itemsets as subsets
                        $C_{k+1} = \text{candidates generated from } L_k$ 
    for each transaction  $t$  in database do
        prove candidates { Increment the count of all candidates in  $C_{k+1}$ 
                           that are contained in  $t$ 
                            $L_{k+1} = \text{candidates in } C_{k+1} \text{ with min\_support}$ 
    return  $\cup_k L_k$ 
```

Generating Candidates (Prune Step)

Step 2: Pruning: $L_{k+1} = \{X \in C_{k+1} \mid \text{support}(X) \geq \text{minSup}\}$

- Naïve: Check support of every itemset in $C_{k+1} \rightarrow$ inefficient for huge C_{k+1}
- Instead, apply Apriori principle first:
Remove candidate $(k+1)$ -itemsets which contain a non-frequent k -subset s , i.e., $s \notin L_k$

```
forall itemsets  $c$  in  $C_{k+1}$  do  
    forall  $k$ -subsets  $s$  of  $c$  do  
        if ( $s$  is not in  $L_k$ ) then delete  $c$  from  $C_{k+1}$ 
```

Example 1

- $L_3 = \{(ACF), (ACG), (AFG), (AFH), (CFG)\}$
- Candidates after the join step: $\{(ACFG), (AFGH)\}$
- In the pruning step: delete (AFGH) because (FGH) $\notin L_3$,
i.e., (FGH) is not a frequent 3-itemset; **also** (AGH) $\notin L_3$
 $\rightarrow C_4 = \{(ACFG)\} \rightarrow$ check the support to generate L_4

Generating Candidates (Join Step)

Requirements for set of all candidate $(k+1)$ -itemsets C_{k+1}

- **Completeness**: Must contain all frequent $(k+1)$ -itemsets (superset property $C_{k+1} \supseteq L_{k+1}$)
- **Selectiveness**: Significantly smaller than the set of all $(k+1)$ -subsets

Step 1: Joining ($C_{k+1} = L_k \bowtie L_k$)


- Suppose the items are sorted by any order (e.g. lexicograph.)
- Consider frequent k -itemsets p and q
- p and q are joined if they share the same first $k-1$ items

insert into C_{k+1}

select $p.i_1, p.i_2, \dots, p.i_{k-1}, p.i_k, q.i_k$

from $L_k : p, L_k : q$

where $p.i_1=q.i_1, \dots, p.i_{k-1}=q.i_{k-1}, p.i_k < q.i_k$



Apriori Algorithm – Full Example

minSup=0.5
database D

TID	items
100	1 3 4 6
200	2 3 5
300	1 2 3 5
400	1 5 6

scan D

itemset	count
{1}	3
{2}	2
{3}	3
{4}	1
{5}	3
{6}	2

itemset	count
{1}	3
{2}	2
{3}	3
{5}	3
{6}	2

$L_1 \bowtie L_1$

itemset
{1 2}
{1 3}
{1 5}
{1 6}
{2 3}
{2 5}
{2 6}
{3 5}
{3 6}
{5 6}

prune C_1

itemset
{1 2}
{1 3}
{1 5}
{1 6}
{2 3}
{2 5}
{2 6}
{3 5}
{3 6}
{5 6}

scan D

itemset	count
{1 2}	1
{1 3}	2
{1 5}	2
{1 6}	2
{2 3}	2
{2 5}	2
{2 6}	0
{3 5}	2
{3 6}	1
{5 6}	1

itemset	count
{1 3}	2
{1 5}	2
{1 6}	2
{2 3}	2
{2 5}	2
{3 5}	2

$L_2 \bowtie L_2$

itemset
{1 3 5}
{1 3 6}
{1 5 6}
{2 3 5}

prune C_2

itemset
{1 3 5}
{1 3 6} x
{1 5 6} x
{2 3 5}

scan D

itemset	count
{1 3 5}	1
{2 3 5}	2

itemset	count
{2 3 5}	2

$L_3 \bowtie L_3$

C_4 is empty

Is Apriori Fast Enough?

– Performance Bottlenecks –

The core of the Apriori algorithm:

- Use frequent $(k - 1)$ -itemsets to generate **candidate** frequent k -itemsets
- Use database scan and pattern matching to collect counts for the candidate itemsets

The bottleneck of *Apriori*: **candidate generation**

- Huge candidate sets:
 - 10^4 frequent 1-itemsets will generate 10^7 candidate 2-itemsets
 - To discover a frequent pattern of size 100, e.g., $\{a_1, a_2, \dots, a_{100}\}$, one needs to generate $2^{100} \approx 10^{30}$ candidates.
 - Multiple scans of database:
 - Needs n or $n+1$ scans, n is the length of the longest pattern
- Is it possible to mine the complete set of frequent itemsets without candidate generation?

Mining Frequent Patterns Without Candidate Generation

Compress a large database into a compact, *Frequent-Pattern tree (FP-tree)* structure

- highly condensed, but complete for frequent pattern mining
- avoid costly database scans

Develop an efficient, FP-tree-based frequent pattern mining method

- A divide-and-conquer methodology: decompose mining tasks into smaller ones
- Avoid candidate generation: sub-database test only!

Idea:

- Compress database into FP-tree, retaining the itemset association information
- Divide the compressed database into conditional databases, each associated with one frequent item and mine each such database separately.

Construct FP-tree from a Transaction DB (1)

Steps for compressing the database into a FP-tree:

1. Scan DB once, find frequent 1-itemsets (single items)
2. Order frequent items in frequency descending order

TID	items bought
100	{f, a, c, d, g, i, m, p}
200	{a, b, c, f, l, m, o}
300	{b, f, h, j, o}
400	{b, c, k, s, p}
500	{a, f, c, e, l, p, m, n}

header table:

minSup=0.5

item	frequency
f	4
c	4
a	3
b	3
m	3
p	3

*sort items in the order of
descending support*

Construct FP-tree (2)

Steps for compressing the database into a FP-tree:

1. Scan DB once, find frequent 1-itemsets (single items)
2. Order frequent items in frequency descending order
3. Scan DB again, construct FP-tree **starting with most frequent item per transaction**

TID	items bought	(ordered) frequent items
100	{f, a, c, d, g, i, m, p}	{f, c, a, m, p}
200	{a, b, c, f, l, m, o}	{f, c, a, b, m}
300	{b, f, h, j, o}	{f, b}
400	{b, c, k, s, p}	{c, b, p}
500	{a, f, c, e, l, p, m, n}	{f, c, a, m, p}

for each transaction only keep its frequent items sorted in descending order of their frequencies

header table:

item	frequency
f	4
c	4
a	3
b	3
m	3
p	3

- for each transaction built a path in the FP-tree:
- If a path with common prefix exists: increment frequency of nodes on this path and append suffix
 - Otherwise: create a new branch


Construct FP-tree (3)

TID	items bought	(ordered) frequent items
100	{f, a, c, d, g, i, m, p}	{f, c, a, m, p}
200	{a, b, c, f, l, m, o}	{f, c, a, b, m}
300	{b, f, h, j, o}	{f, b}
400	{b, c, k, s, p}	{c, b, p}
500	{a, f, c, e, l, p, m, n}	{f, c, a, m, p}

header table:

item	frequency	head
f	4	•
c	4	•
a	3	•
b	3	•
m	3	•
p	3	•

header table
references the
occurrences of the
frequent items in the
FP-tree



Benefits of the FP-tree Structure

- Completeness:
 - never breaks a long pattern of any transaction
 - preserves complete information for frequent pattern mining
- Compactness
 - reduce irrelevant information—infrequent items are gone
 - frequency descending ordering: more frequent items are more likely to be shared
 - never be larger than the original database
(if not count node-links and counts)
 - Experiments demonstrate compression ratios over 100

Mining Frequent Patterns Using FP-tree

- General idea (divide-and-conquer)
 - Recursively grow frequent pattern path using the FP-tree
- Method
 - For each item, construct its **conditional pattern-base (*prefix paths*)**, and then its **conditional FP-tree**
 - Repeat the process on each newly created conditional FP-tree ...
 - ...until the resulting FP-tree is **empty**, or it contains **only one path** (single path will generate all the combinations of its sub-paths, each of which is a frequent pattern)

Major Steps to Mine FP-tree

1. Construct conditional pattern base for each node in the FP-tree
2. Construct conditional FP-tree from each conditional pattern-base
3. Recursively mine conditional FP-trees and grow frequent patterns obtained so far
 - If the conditional FP-tree contains a single path, simply enumerate all the patterns


Major Steps to Mine FP-tree: Conditional Pattern Base

1. Construct conditional pattern base for each node in the FP-tree

- Starting at the frequent header table in the FP-tree
- Traverse the FP-tree by following the link of each frequent item
- Accumulate all of transformed prefix paths of that item to form a conditional pattern base
 - For each item its prefixes are regarded as condition for it being a suffix. These prefixes form the conditional pattern base. The frequency of the prefixes can be read in the node of the item.

header table:

item	frequency	head
<i>f</i>	4	•
<i>c</i>	4	•
<i>a</i>	3	•
<i>b</i>	3	•
<i>m</i>	3	•
<i>p</i>	3	•



conditional pattern base:

item	cond. pattern base
<i>f</i>	{}
<i>c</i>	<i>f</i> :3
<i>a</i>	<i>fc</i> :3
<i>b</i>	<i>fca</i> :1, <i>f</i> :1, <i>c</i> :1
<i>m</i>	<i>fca</i> :2, <i>fcab</i> :1
<i>p</i>	<i>fcam</i> :2, <i>cb</i> :1

Properties of FP-tree for Conditional Pattern Bases

- Node-link property
 - For any frequent item a_i , all the possible frequent patterns that contain a_i can be obtained by following a_i 's node-links, starting from a_i 's head in the FP-tree header
- Prefix path property
 - To calculate the frequent patterns for a node a_i in a path P , only the prefix sub-path of a_i in P needs to be accumulated, and its frequency count should carry the same count as node a_i .

Major Steps to Mine FP-tree: Conditional FP-tree

2. Construct conditional FP-tree from each conditional pattern-base

- The prefix paths of a suffix represent the conditional basis.
→ They can be regarded as transactions of a database.
- Those prefix paths whose support $\geq \text{minSup}$, induce a conditional FP-tree
- For each pattern-base
 - Accumulate the count for each item in the base
 - Construct the FP-tree for the frequent items of the pattern base

conditional pattern base:

item	cond. pattern base
<i>f</i>	<i>{}</i>
<i>c</i>	<i>f:3</i>
<i>a</i>	<i>fc:3</i>
<i>b</i>	<i>fca:1, f:1, c:1</i>
<i>m</i>	<i>fca:2, fcab:1</i>
<i>p</i>	<i>fcam:2, cb:1</i>

e.g. item *m*

item	frequency
<i>f</i>	3 ..
<i>c</i>	3 ..
<i>a</i>	3 ..
<i>b</i>	1 <i>x</i>

m-conditional FP-tree

{}|*m*
|
f:3
|
c:3
|
a:3

Conditional FP-tree

2. Construct conditional FP-tree from each conditional pattern-base

conditional pattern base:

item	cond. pattern base
<i>f</i>	$\{\}$
<i>c</i>	$f:3$
<i>a</i>	$fc:3$
<i>b</i>	$fca:1, f:1, c:1$
<i>m</i>	$fca:2, fcab:1$
<i>p</i>	$fcam:2, cb:1$

$\{\} | f = \{\}$

$\{\} | c$

$f:3$

$\{\} | a$

$f:3$
 $c:3$

$\{\} | b = \{\}$

$\{\} | m$

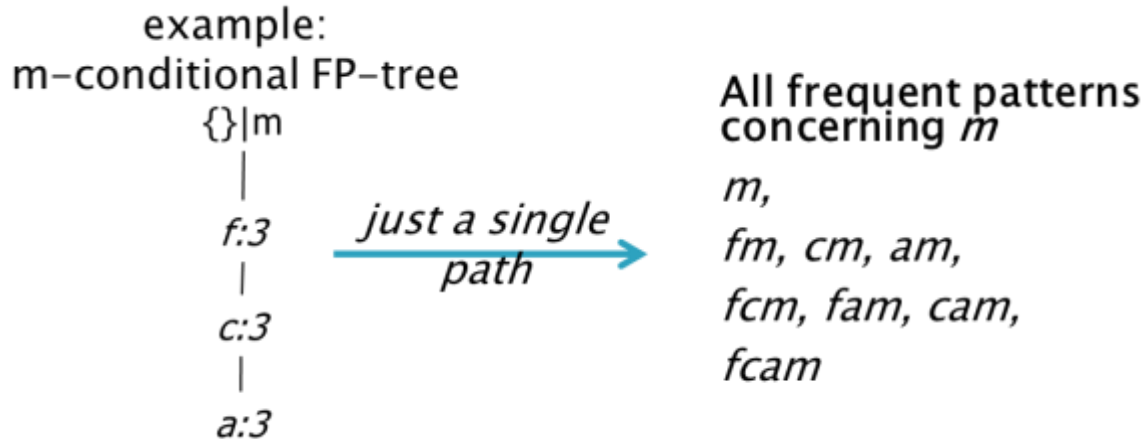
$f:3$
 $c:3$
 $a:3$

$\{\} | p$

$c:3$

Major Steps to Mine FP-tree

3. Recursively mine conditional FP-trees and grow frequent patterns obtained so far
- If the conditional FP-tree contains a single path, simply enumerate all the patterns (enumerate all combinations of sub-paths)



FP-tree: Full Example

database:

TID	items bought	(ordered) frequent items
100	{b, c, f}	{f, b, c}
200	{a, b, c}	{b, c}
300	{d, f}	{f}
400	{b, c, e, f}	{f, b, c}
500	{f, g}	{f}


$minSup=0.4$

header table:

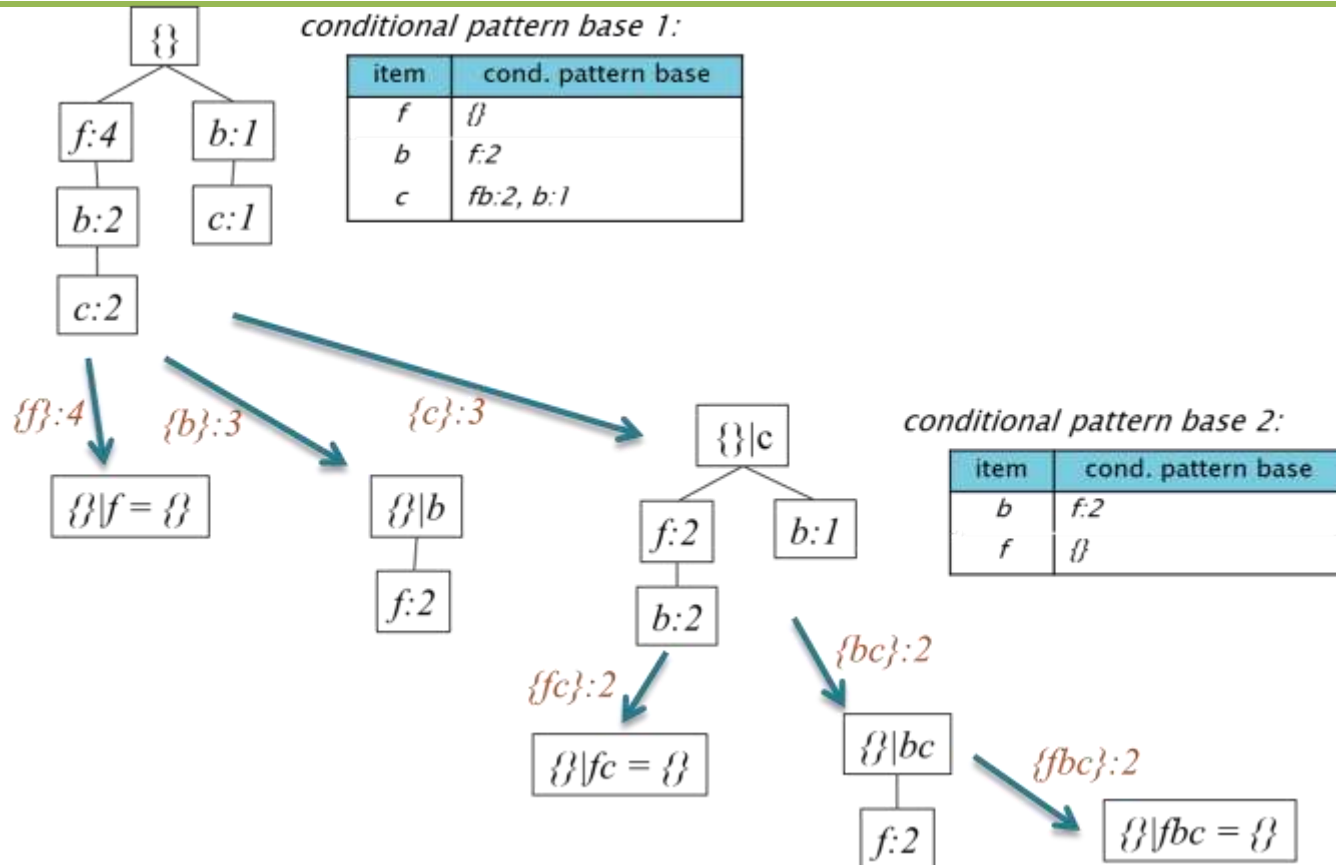
item	frequency	head
f	4	•
b	3	•
c	3	•

conditional pattern base:

item	cond. pattern base
f	{}
b	f:2
c	fb:2, b:1



FP-tree: Full Example



Principles of Frequent Pattern Growth

- Pattern growth property

- Let α be a frequent itemset in DB, B be α 's conditional pattern base, and β be an itemset in B.

Then $\alpha \cup \beta$ is a frequent itemset in DB iff β is frequent in B.

- "abcdef" is a frequent pattern, if and only if

- "abcde" is a frequent pattern, and
- "f" is frequent in the set of transactions containing "abcde"

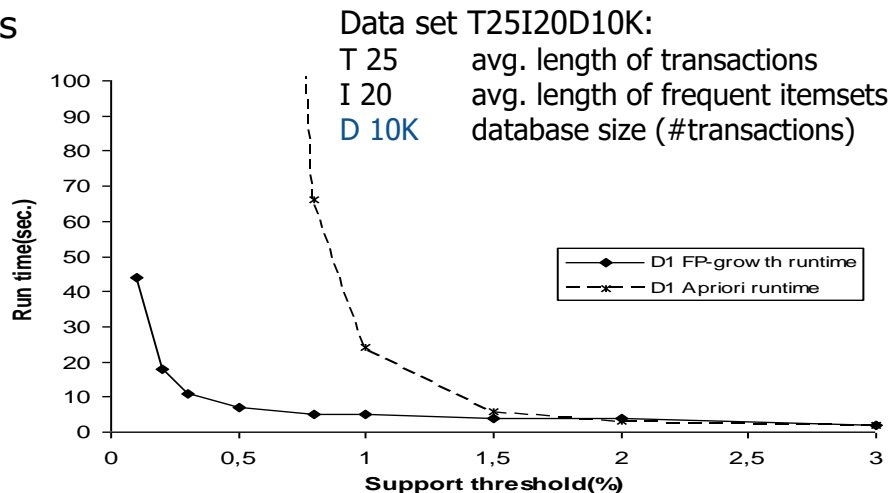
Why Is Frequent Pattern Growth Fast?

Performance study in [Han, Pei&Yin '00] shows

- FP-growth is an order of magnitude faster than Apriori, and is also faster than tree-projection

Reasoning

- No candidate generation, no candidate test
 - Apriori algorithm has to proceed breadth-first
- Use compact data structure
- Eliminate repeated database scan
- Basic operation is counting and FP-tree building



Content Overview

- Introduction
 - Transaction databases, market basket data analysis
- Mining Frequent Itemsets
 - Apriori algorithm, hash trees, FP-tree
- Summerizing Frequent Itemsets
 - Maximal, closed, non-derivable itemsets
- Simple Association Rules
 - Basic notions, rule generation, interestingness measures
- Summary

cf. Zaki's Book!

Maximal & Closed Frequent Itemsets

Big challenge: database contains potentially a huge number of frequent itemsets (especially if minSup is set too low).

- A frequent itemset of length 100 contains $2^{100}-1$ many frequent subsets

Closed frequent itemset:

An itemset X is *closed* in a data set D if there exists no proper super-itemset Y such that $support(X) = support(Y)$ in D .

- The set of closed frequent itemsets contains complete information regarding its corresponding frequent itemsets.

Maximal frequent itemset:

An itemset X is *maximal* in a data set D if there exists no proper super-itemset Y such that $support(Y) \geq minSup$ in D .

- The set of maximal itemsets does not contain the complete support information
- More compact representation

Maximal Frequent Itemsets

Given a binary database $\mathbf{D} \subseteq \mathcal{T} \times \mathcal{I}$, over the tids \mathcal{T} and items \mathcal{I} , let \mathcal{F} denote the set of all frequent itemsets, that is,

$$\mathcal{F} = \{X \mid X \subseteq \mathcal{I} \text{ and } \text{sup}(X) \geq \text{minsup}\}$$

A frequent itemset $X \in \mathcal{F}$ is called *maximal* if it has no frequent supersets. Let \mathcal{M} be the set of all maximal frequent itemsets, given as

$$\mathcal{M} = \{X \mid X \in \mathcal{F} \text{ and } \nexists Y \supset X, \text{ such that } Y \in \mathcal{F}\}$$

The set \mathcal{M} is a condensed representation of the set of all frequent itemset \mathcal{F} , because we can determine whether any itemset X is frequent or not using \mathcal{M} . If there exists a maximal itemset Z such that $X \subseteq Z$, then X must be frequent; otherwise X cannot be frequent.

Example for Maximal Itemsets

Transaction database

Tid	Itemset
1	ABDE
2	BCE
3	ABDE
4	ABCE
5	ABCDE
6	BCD

Frequent itemsets ($minsup = 3$)

sup	Itemsets
6	B
5	E, BE
4	A, C, D, AB, AE, BC, BD, ABE
3	AD, CE, DE, ABD, ADE, BCE, BDE, ABDE

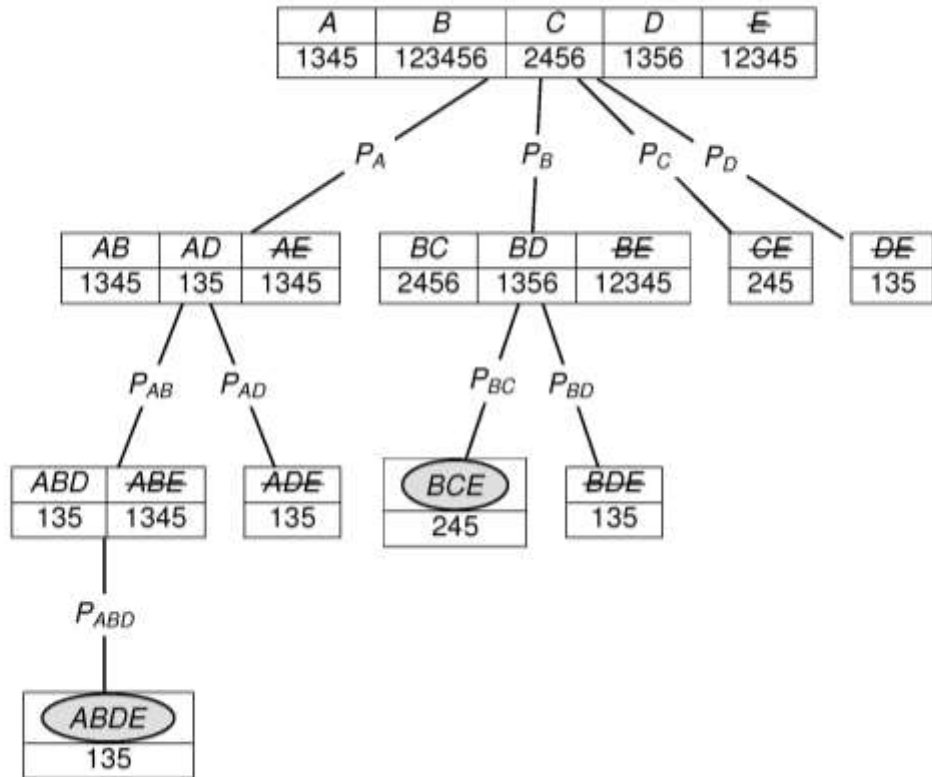
Mining Maximal Frequent Itemsets

Mining maximal itemsets requires additional steps beyond simply determining the frequent itemsets. Assuming that the set of maximal frequent itemsets is initially empty, that is, $\mathcal{M} = \emptyset$, each time we generate a new frequent itemset X , we have to perform the following maximality checks

- **Subset Check:** $\nexists Y \in \mathcal{M}$, such that $X \subset Y$. If such a Y exists, then clearly X is not maximal. Otherwise, we add X to \mathcal{M} , as a potentially maximal itemset.
- **Superset Check:** $\nexists Y \in \mathcal{M}$, such that $Y \subset X$. If such a Y exists, then Y cannot be maximal, and we have to remove it from \mathcal{M} .

GenMax Algorithm

```
// Initial Call:  $\mathcal{M} \leftarrow \emptyset$ ,  
                $P \leftarrow \{\langle i, \mathbf{t}(i) \rangle \mid i \in \mathcal{I}, \text{sup}(i) \geq \text{minsup}\}$   
GENMAX ( $P, \text{minsup}, \mathcal{M}$ ):  
1  $Y \leftarrow \bigcup X_i$   
2 if  $\exists Z \in \mathcal{M}$ , such that  $Y \subseteq Z$  then  
3   return // prune entire branch  
4 foreach  $\langle X_i, \mathbf{t}(X_i) \rangle \in P$  do  
5    $P_i \leftarrow \emptyset$   
6   foreach  $\langle X_j, \mathbf{t}(X_j) \rangle \in P$ , with  $j > i$  do  
7      $X_{ij} \leftarrow X_i \cup X_j$   
8      $\mathbf{t}(X_{ij}) = \mathbf{t}(X_i) \cap \mathbf{t}(X_j)$   
9     if  $\text{sup}(X_{ij}) \geq \text{minsup}$  then  $P_i \leftarrow P_i \cup \{\langle X_{ij}, \mathbf{t}(X_{ij}) \rangle\}$   
10  if  $P_i \neq \emptyset$  then GENMAX ( $P_i, \text{minsup}, \mathcal{M}$ )  
11  else if  $\nexists Z \in \mathcal{M}, X_i \subseteq Z$  then  
12     $\mathcal{M} = \mathcal{M} \cup X_i$  // add  $X_i$  to maximal set
```



Content Overview

- Introduction
 - Transaction databases, market basket data analysis
- Mining Frequent Itemsets
 - Apriori algorithm, hash trees, FP-tree
- Summerizing Frequent Itemsets
 - Maximal, closed, non-derivable itemsets
- Simple Association Rules
 - Basic notions, rule generation, interestingness measures
- Summary

Association Rules: Basic Notions

Items $I = \{i_1, i_2, \dots, i_m\}$: a set of literals (denoting items)

Itemset X : Set of items $X \subseteq I$

Database D : Set of *transactions* T , each transaction is a set of items $T \subseteq I$

Transaction T *contains* an itemset X : $X \subseteq T$

The *support* of an itemset X is defined as: $\text{support}(X) = |\{T \in D | X \subseteq T\}|$

Frequent itemset: an itemset X is called frequent iff $\text{support}(X) \geq \text{minSup}$

Association rule: An association rule is an implication in the form $A \Rightarrow B$ where $A \subset I$ and $B \subset I$ are two itemsets with $A \cap B = \emptyset$

Note: simply enumerating all possible association rules is not meaningful!
→ What are the interesting association rules w.r.t. D ?

Interestingness of Association Rules

Interestingness of an association rule:

Quantify the interestingness of an association rule with respect to a transaction database D :

- Support: frequency (probability) of the entire rule with respect to D

$$\text{support}(A \Rightarrow B) = P(A \cup B) = \frac{|\{T \in D | A \cup B \subseteq T\}|}{|D|} = \text{support}(A \cup B)$$

“probability that a transaction in D contains the itemset $A \cup B$ ”

- Confidence: indicates the strength of implication in the rule


$$\text{confidence}(A \Rightarrow B) = P(B|A) = \frac{|\{T \in D | A \cup B \subseteq T\}|}{|\{T \in D | A \subseteq T\}|} = \frac{\text{support}(A \cup B)}{\text{support}(A)}$$

“conditional probability that a transaction in D containing the itemset A also contains itemset B ”

- Rule form: “*Body* \Rightarrow *Head* [*support*, *confidence*]”

Association rule examples:


- buys diapers \Rightarrow buys beers [0.5%, 60%]
- major in CS \wedge takes DB \Rightarrow avg. grade A [1%, 75%]



Mining of Association Rules


Given a database D , determine all association rules having a $support \geq minSup$ and a $confidence \geq minConf$ (so-called *strong association rules*).

Key steps of mining association rules:

- e.g.  Apriori, FP-growth
- 1) Find *frequent itemsets*, i.e., itemsets that have at least $support = minSup$
 - 2) Use the frequent itemsets to generate association rules
 - n frequent items yield $2^n - 2$ association rules

Example

transaction ID	items
2000	A, B, C
1000	A, C
4000	A, D
5000	B, E, F

$minSup = 50\%$


frequent items	support
{A}	75%
{B}	50%
{C}	50%
{A, C}	50%

For $minSup = 50\%$ and $minConf = 50\%$ we obtain 2 rules:

- $A \Rightarrow C$: $support(A \Rightarrow C) = 50\%$, $confidence(A \Rightarrow C) = 66.67\%$
- $C \Rightarrow A$: $support(C \Rightarrow A) = 50\%$, $confidence(C \Rightarrow A) = 100\%$

Generating Rules

For each frequent itemset X

- For each nonempty subset Y of X , form a rule $Y \Rightarrow (X - Y)$
- Delete those rules that do not have minimum confidence

Computation of the confidence of a rule $Y \Rightarrow (X - Y)$

$$\text{confidence}(Y \Rightarrow (X - Y)) = \frac{\text{support}(X)}{\text{support}(Y)}$$

Store the frequent itemsets and their support in a hash table in main memory \rightarrow no additional database access

Example: $X = \{A, B, C\}$, $\text{minConf} = 60\%$

- $\text{conf}(A \Rightarrow B, C) = 2/2$; ✓ $\text{conf}(B, C \Rightarrow A) = 1/2$ ✗
- $\text{conf}(B \Rightarrow A, C) = 2/4$; ✗ $\text{conf}(A, C \Rightarrow B) = 1$ ✓
- $\text{conf}(C \Rightarrow A, B) = 2/5$; ✗ $\text{conf}(A, B \Rightarrow C) = 2/3$ ✓

itemset	support
{A}	2
{B}	4
{C}	5
{A, B}	3
{A, C}	2
{B, C}	4
{A, B, C}	2

Interestingness Measurements

Objective measures

- Two popular measurements:
- **support** and
- **confidence**

Subjective measures

[Silberschatz & Tuzhilin, KDD95]

- A rule (pattern) is interesting if it is
- **unexpected** (surprising to the user) and/or
- **actionable** (the user can do something with it)

Criticism to Support and Confidence

Example 1 [Aggarwal & Yu, PODS98]

Among 5000 students

- 3000 play basketball (=60%)
- 3750 eat cereal (=75%)
- 2000 both play basket ball and eat cereal (=40%)

Rule *play basketball* \Rightarrow *eat cereal* [40%, 66.7%] is **misleading** because the overall percentage of students eating cereal is 75% which is higher than 66.7%

Rule *play basketball* \Rightarrow *not eat cereal* [20%, 33.3%] is far **more accurate**, although with lower support and confidence

Observation: *play basketball* and *eat cereal* are **negatively correlated**

Other Interestingness Measures: Correlation

Lift is a simple correlation measure between two items A and B :

$$\text{corr}_{A,B} = \frac{P(A \cup B)}{P(A)P(B)} = \frac{P(A|B)}{P(A)} = \frac{P(B|A)}{P(B)}$$

! The two rules $A \Rightarrow B$ and $B \Rightarrow A$ have the same correlation coefficient

take both $P(A)$ and $P(B)$ in consideration

$\text{corr}_{A,B} > 1$ the two items A and B are positively correlated

$\text{corr}_{A,B} = 1$ there is no correlation between the two items A and B

$\text{corr}_{A,B} < 1$ the two items A and B are negatively correlated

Summary

Mining Frequent Itemsets

- Apriori algorithm, FP-tree

Summarizing Frequent Itemsets

- Maximal itemsets

Simple Association Rules

- Basic notions, rule generation, interestingness measures